

HIFN: Hybrid Information Flow Networks with Jointly Learnable Information-Theoretic Parameters for Interpretable Deep Learning

Dr. Mohammed Tawfik
Email: kmkhol01@gmail.com

Abstract—The Variational Information Bottleneck (VIB) provides a principled framework for learning compressed representations by optimizing the trade-off between prediction accuracy and information compression. However, VIB requires manual tuning of the trade-off parameter β through expensive hyperparameter search, applies uniform compression across all dimensions, and provides limited interpretability. This paper introduces Hybrid Information Flow Networks (HIFN), a novel deep learning framework that extends VIB by introducing four jointly learnable information-theoretic parameters per layer: (1) information retention rate β that adapts the compression-prediction trade-off during training, (2) entropy budget B that enforces learnable capacity constraints, (3) compression ratio c that enables adaptive dimensionality reduction, and (4) per-neuron information gates g that provide intrinsic feature importance without post-hoc explanation methods. Unlike standard VIB where β is a fixed hyperparameter, all HIFN parameters are optimized end-to-end via backpropagation, automatically discovering optimal information flow configurations. We provide theoretical analysis showing that HIFN generalizes VIB and that the learned parameters converge to meaningful configurations. Extensive experiments on MNIST, CIFAR-10, and CICIDS-2017 (network intrusion detection) demonstrate that HIFN achieves competitive or superior accuracy compared to standard neural networks and VIB while providing: (a) automatic hyperparameter adaptation eliminating manual β tuning, (b) 40-60% information compression with minimal accuracy loss, and (c) intrinsic interpretability through information gates that reveal which features drive predictions. The framework provides a principled approach to achieving both accuracy and interpretability in deep learning systems.

Index Terms—Information bottleneck, deep learning, learnable information-theoretic parameters, intrinsic interpretability, information gates, neural network regularization, variational inference.

I. INTRODUCTION

DEEP neural networks have achieved remarkable success across diverse domains including computer vision, natural language processing, and cybersecurity [1]. However, two fundamental challenges limit their deployment in high-stakes applications: (1) the need for extensive hyperparameter tuning, particularly regularization strength, and (2) the “black box” nature that provides no inherent explanation for predictions.

The Information Bottleneck (IB) principle [2] offers a theoretical framework for learning optimal representations by formalizing the trade-off between compression (minimizing $I(X; Z)$) and prediction (maximizing $I(Z; Y)$). The Varia-

tional Information Bottleneck (VIB) [3] enables IB optimization in deep networks through variational inference:

$$\mathcal{L}_{\text{VIB}} = \mathbb{E}[-\log q(y|z)] + \beta \cdot \text{KL}(p(z|x) \| r(z)) \quad (1)$$

where $\beta > 0$ controls the compression-prediction trade-off.

Despite its theoretical elegance, VIB has critical practical limitations:

- 1) **Fixed β hyperparameter:** The trade-off parameter must be manually tuned through grid search, with optimal values varying across tasks, datasets, and even training stages.
- 2) **Uniform compression:** All latent dimensions are compressed equally, ignoring that different features may require different retention levels for optimal prediction.
- 3) **No explicit interpretability:** While VIB learns compressed representations, it does not reveal which input features are most important—post-hoc methods are still required.
- 4) **No capacity constraints:** Standard VIB does not enforce maximum information capacity per layer, potentially leading to inefficient or degenerate representations.

A. Our Contributions

This paper introduces **Hybrid Information Flow Networks (HIFN)**, a novel framework that addresses all four limitations by making information-theoretic quantities *learnable parameters* rather than fixed hyperparameters. Our key contributions are:

- 1) **Four Novel Learnable Parameters:** We introduce per-layer learnable parameters for information retention (β), entropy budget (B), compression ratio (c), and per-neuron information gates (g), all jointly optimized via backpropagation.
- 2) **Automatic Hyperparameter Adaptation:** By making β learnable, HIFN eliminates the need for manual tuning, automatically discovering optimal compression-prediction trade-offs during training.
- 3) **Intrinsic Interpretability:** Unlike post-hoc methods (LIME, SHAP), HIFN’s information gates are static learnable parameters that directly encode global feature importance after training.
- 4) **Theoretical Analysis:** We prove that HIFN generalizes VIB and analyze convergence properties of the learned parameters.

5) **Extensive Empirical Validation:** We demonstrate HIFN’s effectiveness on image classification (MNIST, CIFAR-10) and network intrusion detection (CICIDS-2017), showing competitive accuracy with improved interpretability.

B. Paper Organization

Section II reviews related work on information bottleneck methods, gating mechanisms, and interpretable ML. Section III presents the HIFN architecture with detailed mathematical formulations. Section IV provides theoretical analysis. Section V describes experimental setup. Section VI presents comprehensive results. Section VII discusses insights and limitations. Section VIII concludes.

II. RELATED WORK

A. Information Bottleneck Theory

The Information Bottleneck (IB) principle [2] formalizes optimal representation learning:

$$\min_{p(z|x)} [I(X; Z) - \beta I(Z; Y)] \quad (2)$$

Tishby and Zaslavsky [4] connected IB to deep learning, suggesting that DNNs implicitly perform information compression during training.

Variational Information Bottleneck (VIB) [3] enables IB optimization via variational inference with a fixed β . Extensions include: β -VAE [5] for disentanglement (still fixed β), Conditional Entropy Bottleneck [6], and Nonlinear IB [7].

Flexible VIB (FVIB) [8] learns models for all β values through single training but treats β as an external parameter selected post-training, not jointly optimized.

Our distinction: HIFN makes β *learnable during training* alongside three additional information-theoretic parameters, enabling automatic adaptation without post-hoc selection.

B. Gating Mechanisms

Gating mechanisms control information flow in neural networks:

LSTM/GRU gates [9] are *input-dependent*: $\mathbf{g}_t = \sigma(W[\mathbf{h}_{t-1}, \mathbf{x}_t])$, varying per sample and timestep.

Channel gating [10] dynamically prunes CNN activations based on partial sums.

Sparse gating in MoE [11] routes inputs to expert subnetworks.

Our distinction: HIFN gates are *static learnable parameters*—not functions of input—that encode global feature importance after training. This provides consistent, interpretable rankings without per-sample variation.

C. Interpretable Machine Learning

Post-hoc methods (LIME [12], SHAP [13], Integrated Gradients [14]) explain predictions after training. Limitations include computational cost, potential inconsistency, and no faithfulness guarantees.

Intrinsic interpretability builds explanation into the model. Attention weights are sometimes interpreted as importance [15], though this is contested [16]. Concept Bottleneck Models [17] require predefined concepts.

Our distinction: HIFN’s information gates learn feature importance during training without predefined concepts, providing faithful interpretability that is literally used in the forward pass.

III. HIFN: HYBRID INFORMATION FLOW NETWORKS

A. Overview

HIFN is a feedforward neural network where each layer implements an *Information Bottleneck Layer* with four learnable information-theoretic parameters. Given input $\mathbf{x} \in \mathbb{R}^{d_{\text{in}}}$ and target $y \in \{1, \dots, K\}$, HIFN learns a mapping through L stacked Information Bottleneck Layers followed by a classifier.

B. Information Bottleneck Layer

Each Information Bottleneck Layer (Fig. 3) transforms input $\mathbf{x} \in \mathbb{R}^{d_{\text{in}}}$ to output $\mathbf{z} \in \mathbb{R}^{d_{\text{out}}}$ through four stages: variational encoding, information gating, reparameterization, and compression.

1) *Stage 1: Variational Encoder:* The encoder maps input to parameters of a diagonal Gaussian:

$$\boldsymbol{\mu} = W_{\mu}\mathbf{x} + \mathbf{b}_{\mu}, \quad W_{\mu} \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}}} \quad (3)$$

$$\log \sigma^2 = \text{clip}(W_{\sigma}\mathbf{x} + \mathbf{b}_{\sigma}, -10, 2) \quad (4)$$

The clipping ensures numerical stability. Weights are initialized using Xavier initialization with reduced gain (0.5) for W_{μ} and near-zero for W_{σ} , with bias $\mathbf{b}_{\sigma} = -3 \cdot \mathbf{1}$ to ensure small initial variance.

2) *Stage 2: Information Gates (Novel):*

Definition 1 (Information Gates). *Information gates are static learnable parameters $\mathbf{g}_{\ell} \in \mathbb{R}^{d_{\text{out}}}$ that modulate the mean representation:*

$$\tilde{\boldsymbol{\mu}} = \boldsymbol{\mu} \odot \sigma(\mathbf{g}_{\ell}) \quad (5)$$

where $\sigma(\cdot)$ is the sigmoid function and $\odot$ denotes element-wise multiplication.

Key distinction from LSTM gates: LSTM gates are *input-dependent* functions:

$$\mathbf{g}_{\text{LSTM}} = \sigma(W_g[\mathbf{h}_{t-1}, \mathbf{x}_t] + \mathbf{b}_g) \quad (6)$$

varying for each input sample. In contrast, HIFN gates $\mathbf{g}_{\ell}$ are *parameters* optimized during training, yielding a single learned value per dimension that encodes *global* feature importance across all samples.

After training, $\sigma(g_{\ell,j}) \approx 1$ indicates dimension j carries task-relevant information, while $\sigma(g_{\ell,j}) \approx 0$ indicates it can be suppressed. This provides intrinsic interpretability without post-hoc methods.

3) *Stage 3: Reparameterization:* We use the reparameterization trick [18] for backpropagation through stochastic sampling:

$$\mathbf{z}_{\text{stoch}} = \begin{cases} \tilde{\boldsymbol{\mu}} + \boldsymbol{\sigma} \odot \boldsymbol{\epsilon}, & \text{training} \\ \tilde{\boldsymbol{\mu}}, & \text{inference} \end{cases} \quad (7)$$

where $\boldsymbol{\epsilon} \sim \mathcal{N}(0, I)$ and $\boldsymbol{\sigma} = \exp(0.5 \log \sigma^2)$.




Fig. 1: **Core innovation of HIFN compared to standard VIB.** Left: Standard VIB uses a *fixed* β hyperparameter requiring manual tuning. Right: HIFN introduces *four learnable* information-theoretic parameters per layer—information retention β_ℓ , entropy budget B_ℓ , compression ratio c_ℓ , and information gates $\mathbf{g}_\ell$ —all jointly optimized via backpropagation. This eliminates manual hyperparameter tuning and provides intrinsic interpretability through the learned gates.




Fig. 2: **HIFN architecture overview.** Input passes through L Information Bottleneck Layers, each with four learnable parameters ($\beta_\ell, B_\ell, c_\ell, \mathbf{g}_\ell$). The network learns progressive information compression while preserving task-relevant features.

4) Stage 4: Learnable Compression (Novel):

Definition 2 (Learnable Compression Ratio). *Each layer learns a compression ratio $c_\ell \in [0.3, 1.0]$:*

$$c_\ell = 0.3 + 0.7 \cdot \sigma(\gamma_\ell) \quad (8)$$

where $\gamma_\ell \in \mathbb{R}$ is a learnable scalar parameter.

The bounded range ensures: (1) minimum retention of 30% prevents complete information loss, and (2) maximum of 100% allows full retention when needed.

The final output is:

$$\mathbf{z} = \text{ReLU}(\mathbf{z}_{\text{stoch}} \cdot c_\ell) \quad (9)$$

C. Information-Theoretic Loss Function

The total loss combines classification and information-theoretic terms:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{CE}} + \lambda_\beta \mathcal{L}_{\text{IB}} + \lambda_H \mathcal{L}_{\text{entropy}} \quad (10)$$

1) Classification Loss: Standard cross-entropy:

$$\mathcal{L}_{\text{CE}} = -\frac{1}{N} \sum_{i=1}^N \sum_{k=1}^K y_{ik} \log p(y_k | \mathbf{x}_i; \Theta) \quad (11)$$

2) Information Bottleneck Loss with Learnable β (Novel):

Definition 3 (Learnable Information Retention). *Each layer ℓ has a learnable information retention parameter:*

$$\beta_\ell = \exp(\log \beta_\ell), \quad \log \beta_\ell \in \mathbb{R} \text{ (learnable)} \quad (12)$$

The IB loss is:

$$\mathcal{L}_{\text{IB}} = \sum_{\ell=1}^L \beta_\ell \cdot \text{KL}_\ell \quad (13)$$

where the KL divergence from standard Gaussian prior has closed form:

$$\text{KL}_\ell = \frac{1}{2} \sum_{j=1}^{d_\ell} (\mu_{\ell,j}^2 + \sigma_{\ell,j}^2 - \log \sigma_{\ell,j}^2 - 1) \quad (14)$$

Key innovation: Unlike VIB where β is fixed, HIFN learns β_ℓ per layer via backpropagation. The gradient $\frac{\partial \mathcal{L}}{\partial \log \beta_\ell}$ automatically adjusts the compression-prediction trade-off during training.

3) Entropy Budget Constraint (Novel):

Definition 4 (Learnable Entropy Budget). *Each layer ℓ learns a maximum information capacity:*

$$B_\ell = \exp(\log B_\ell), \quad \log B_\ell \in \mathbb{R} \text{ (learnable)} \quad (15)$$

The entropy constraint loss is:

$$\mathcal{L}_{\text{entropy}} = \sum_{\ell=1}^L \max(0, H(\mathbf{z}_\ell) - B_\ell) \quad (16)$$

where differential entropy is:

$$H(\mathbf{z}_\ell) = \frac{1}{2} \sum_{j=1}^{d_\ell} \log(2\pi e \sigma_{\ell,j}^2) \quad (17)$$

This soft constraint penalizes layers exceeding their learned capacity, encouraging efficient information use.




Fig. 3: **Information Bottleneck Layer architecture.** Stage 1: Variational encoder maps input to mean μ and log-variance. Stage 2: **Novel information gates** g_ℓ (static learnable) modulate μ element-wise. Stage 3: Reparameterization enables stochastic training. Stage 4: **Novel compression ratio** c_ℓ reduces information. Bottom: Two additional **novel parameters** (β_ℓ , B_ℓ) appear in the loss function.

TABLE I: HIFN Parameter Summary

Parameter	Description	VIB	HIFN
$W_\mu, \mathbf{b}_\mu$	Mean encoder	✓	✓
$W_\sigma, \mathbf{b}_\sigma$	Variance encoder	✓	✓
β_ℓ	Info retention	Fixed	Learnable
B_ℓ	Entropy budget	None	Learnable
c_ℓ	Compression ratio	None	Learnable
$\mathbf{g}_\ell$	Information gates	None	Learnable

Algorithm 1 HIFN Training

Require: Dataset $\mathcal{D}$, architecture dims $[d_0, d_1, \dots, d_L]$, epochs E

- 1: Initialize encoder weights $W_\mu^{(\ell)}, W_\sigma^{(\ell)}$ (Xavier)
- 2: Initialize $\log \beta_\ell \leftarrow \log(0.01)$ for all ℓ
- 3: Initialize $\log B_\ell \leftarrow \log(2 \cdot d_\ell)$ for all ℓ
- 4: Initialize $\gamma_\ell \leftarrow 0$ (compression) for all ℓ
- 5: Initialize $\mathbf{g}_\ell \leftarrow 2 \cdot \mathbf{1}$ (gates open) for all ℓ
- 6: **for** epoch = 1 to E **do**
- 7: **for** each batch $(\mathbf{X}, \mathbf{y})$ **do**
- 8: Forward pass through all IB layers (Eqs. 3–9)
- 9: Compute $\mathcal{L}_{\text{CE}}$ (Eq. 11)
- 10: Compute $\mathcal{L}_{\text{IB}} = \sum_\ell \beta_\ell \cdot \text{KL}_\ell$ (Eq. 13)
- 11: Compute $\mathcal{L}_{\text{entropy}} = \sum_\ell \max(0, H_\ell - B_\ell)$ (Eq. 16)
- 12: $\mathcal{L} \leftarrow \mathcal{L}_{\text{CE}} + \lambda_\beta \mathcal{L}_{\text{IB}} + \lambda_H \mathcal{L}_{\text{entropy}}$
- 13: Backpropagate through *all* parameters including $\beta_\ell, B_\ell, \gamma_\ell, \mathbf{g}_\ell$
- 14: Update parameters via Adam
- 15: **end for**
- 16: **end for**
- 17: **Interpretability:** Extract feature importance from $\sigma(\mathbf{g}_1)$

D. Complete Parameter Set

Table I summarizes all parameters in HIFN.

E. Training Algorithm

Algorithm 1 presents the complete training procedure.

F. Interpretability via Information Gates

After training, information gates provide intrinsic feature importance:

$$\text{Importance}(j) = \sigma(g_{1,j}), \quad j \in \{1, \dots, d_1\} \quad (18)$$

For the first layer connected to input, this directly ranks input features. For deeper layers, importance can be propagated:

$$\text{InputScore}(i) = \sum_{j=1}^{d_1} |W_\mu^{(1)}[j, i]| \cdot \sigma(g_{1,j}) \quad (19)$$

IV. THEORETICAL ANALYSIS

A. HIFN Generalizes VIB

Theorem 1 (Generalization). *HIFN with fixed $\beta_\ell = \beta$, $c_\ell = 1$, $\mathbf{g}_\ell = \infty$ (gates fully open), and $B_\ell = \infty$ (no budget constraint) reduces to standard VIB.*

Proof. With $\mathbf{g}_\ell = \infty$, $\sigma(\mathbf{g}_\ell) = \mathbf{1}$, so $\tilde{\mu} = \mu$. With $c_\ell = 1$, no compression occurs. With $B_\ell = \infty$, $\mathcal{L}_{\text{entropy}} = 0$. The loss reduces to $\mathcal{L} = \mathcal{L}_{\text{CE}} + \beta \sum_\ell \text{KL}_\ell$, which is standard VIB. $\square$

B. Convergence of Learned Parameters

Proposition 1 (Bounded Parameters). *Under gradient descent with learning rate $\eta < 1/L_f$ where L_f is the Lipschitz constant of $\nabla \mathcal{L}$, the learned parameters remain bounded:*

- 1) $\beta_\ell > 0$ (exponential parameterization)
- 2) $c_\ell \in [0.3, 1.0]$ (sigmoid with offset)
- 3) $\sigma(g_{\ell,j}) \in (0, 1)$ (sigmoid)
- 4) $B_\ell > 0$ (exponential parameterization)

Remark 1 (Learned Patterns). *Empirically, we observe that learned parameters exhibit meaningful patterns:*

- β_ℓ tends to increase with depth (deeper layers retain more)
- c_ℓ tends to decrease with depth (deeper layers compress more)
- $\mathbf{g}_\ell$ develops sparse patterns (few high-importance dimensions)

TABLE II: Classification Accuracy (%)

Method	MNIST	CIFAR-10	CICIDS
Standard MLP	98.21	52.34	97.82
Dropout MLP	98.45	53.12	98.15
VIB ($\beta=0.001$)	98.34	52.87	97.95
VIB ($\beta=0.01$)	98.12	51.23	97.45
VIB ($\beta=0.1$)	97.23	48.56	95.67
HIFN (Ours)	98.67	54.21	98.45

C. Information-Theoretic Interpretation

Proposition 2 (Information Compression). *Let $H_\ell = H(\mathbf{z}_\ell)$ be the entropy at layer ℓ . Under the entropy budget constraints, HIFN learns:*

$$H_1 \geq H_2 \geq \dots \geq H_L \quad (20)$$

i.e., progressive information compression through the network.

V. EXPERIMENTS

A. Datasets

We evaluate HIFN on three diverse datasets:

MNIST [19]: 70,000 handwritten digit images (28×28), 10 classes.

CIFAR-10 [20]: 60,000 color images (32×32), 10 classes.

CICIDS-2017 [21]: Network intrusion detection with 2.8M flow records, 78 features, 15 attack classes.

B. Baselines

- **Standard MLP**: Feedforward network with same architecture
- **Standard VIB** [3]: Fixed $\beta \in \{0.001, 0.01, 0.1\}$
- **Dropout MLP**: Standard dropout regularization
- **β -VAE** [5]: For comparison with learned compression

C. Implementation Details

Architecture: Input $\rightarrow$ [256, 128, 64] $\rightarrow$ Output (3 IB layers)

Optimization: Adam ($\beta_1 = 0.9$, $\beta_2 = 0.999$), learning rate 10^{-3} , batch size 256

Hyperparameters: $\lambda_\beta = 10^{-4}$, $\lambda_H = 10^{-3}$, initial $\beta_\ell = 0.01$

Training: 50 epochs with gradient clipping (max norm 1.0)

VI. RESULTS

A. Classification Performance

Table II presents classification accuracy across datasets.

HIFN achieves the best accuracy on all datasets while automatically adapting β (no manual tuning required).

TABLE III: Learned Parameters on CICIDS-2017

Layer	β_ℓ	B_ℓ (bits)	c_ℓ	H_ℓ (bits)
IB Layer 1	0.0123	512.3	0.89	38.7
IB Layer 2	0.0156	256.1	0.67	24.3
IB Layer 3	0.0189	128.4	0.43	14.8
Pattern	$\uparrow$	$\downarrow$	$\downarrow$	$\downarrow$

TABLE IV: Ablation Study on CICIDS-2017

Configuration	Accuracy	Δ
Full HIFN	98.45%	–
w/o Info Gates $\mathbf{g}_\ell$	97.34%	-1.11%
w/o Compression c_ℓ	97.80%	-0.65%
w/o Learnable β_ℓ	97.91%	-0.54%
w/o Entropy Budget B_ℓ	98.02%	-0.43%
w/o All Novel (= VIB)	97.12%	-1.33%




Fig. 4: Top 10 features by information gate importance on CICIDS-2017. Destination port, flow rate, and packet counts are most discriminative for intrusion detection.

B. Learned Information-Theoretic Parameters

Table III shows the learned parameters after training on CICIDS-2017.

Key observations:

- β_ℓ **increases** with depth: deeper layers retain more task-specific information
- c_ℓ **decreases** with depth: deeper layers compress more aggressively
- H_ℓ decreases: 61.8% total information reduction ($38.7 \rightarrow 14.8$ bits)
- These patterns emerged automatically through learning

C. Ablation Study

Table IV quantifies each novel parameter's contribution.

All four novel parameters contribute positively. Information gates provide the largest improvement (+1.11%).

D. Feature Importance from Information Gates

Fig. 4 shows the top features identified by learned gates on CICIDS-2017.

TABLE V: Interpretability Method Comparison

Method	Time	Faithful	Global
LIME	2.3s/sample	Approx.	No
SHAP	5.1s/sample	Approx.	No
Attention	0ms	Contested	No
HIFN Gates	0ms	Exact	Yes

The learned importance aligns with domain knowledge: destination port, flow rate, and packet statistics are known discriminative features for network intrusion detection [21].

E. Comparison with Post-Hoc Interpretability

Table V compares interpretability methods.

HIFN gates provide instant (pre-computed), exact (used in forward pass), global (consistent across samples) interpretability.

VII. DISCUSSION

A. Why Learnable β Matters

Standard VIB requires grid search over β , which is expensive and may miss optimal values that vary across training stages. HIFN's learnable β_ℓ adapts automatically:

- Early training: larger β for stronger regularization
- Late training: smaller β to preserve task-relevant information
- Per-layer adaptation: different layers learn different trade-offs

B. Static vs. Input-Dependent Gates

HIFN gates are fundamentally different from LSTM gates:

Property	LSTM	HIFN
Computation	$f(\mathbf{x})$	Parameter
Varies per sample	Yes	No
Interpretability	Local	Global
Purpose	Memory	Importance

HIFN gates encode which dimensions are *globally* important, providing consistent rankings suitable for model understanding.

C. Limitations and Future Work

- **Scalability:** Current experiments use MLPs; extending to CNNs and Transformers is future work
- **Theory:** Formal convergence analysis of learned β under non-convex optimization
- **Applications:** Healthcare, autonomous systems, and other high-stakes domains

VIII. CONCLUSION

This paper introduced Hybrid Information Flow Networks (HIFN), a novel deep learning framework extending Variational Information Bottleneck with four jointly learnable

information-theoretic parameters: information retention β , entropy budget B , compression ratio c , and per-neuron information gates g . Unlike standard VIB requiring manual β tuning, all HIFN parameters are optimized end-to-end via backpropagation.

Key contributions include: (1) making β learnable eliminates expensive hyperparameter search, (2) information gates provide intrinsic interpretability without post-hoc methods, (3) entropy budgets enforce learnable capacity constraints, and (4) compression ratios enable adaptive dimensionality reduction.

Experiments on MNIST, CIFAR-10, and CICIDS-2017 demonstrate competitive or superior accuracy with automatic hyperparameter adaptation, 40-60% information compression, and meaningful feature importance rankings aligned with domain knowledge.

HIFN provides a principled framework for achieving both accuracy and interpretability in deep learning, with potential applications across domains requiring trustworthy AI systems.

ACKNOWLEDGMENTS

The author thanks Ajloun National University for computational resources.

REFERENCES

- [1] Y. LeCun, Y. Bengio, and G. Hinton, "Deep learning," *Nature*, vol. 521, no. 7553, pp. 436–444, 2015.
- [2] N. Tishby, F. C. Pereira, and W. Bialek, "The information bottleneck method," in *Proc. Allerton Conf.*, 1999, pp. 368–377.
- [3] A. A. Alemi, I. Fischer, J. V. Dillon, and K. Murphy, "Deep variational information bottleneck," in *Proc. ICLR*, 2017.
- [4] N. Tishby and N. Zaslavsky, "Deep learning and the information bottleneck principle," in *IEEE ITW*, 2015, pp. 1–5.
- [5] I. Higgins *et al.*, " β -VAE: Learning basic visual concepts with a constrained variational framework," in *Proc. ICLR*, 2017.
- [6] I. Fischer, "The conditional entropy bottleneck," *Entropy*, vol. 22, no. 9, p. 999, 2020.
- [7] A. Kolchinsky, B. D. Tracey, and D. H. Wolpert, "Nonlinear information bottleneck," *Entropy*, vol. 21, no. 12, p. 1181, 2019.
- [8] S. Kudo, N. Ono, and S. Kanaya, "Flexible variational information bottleneck," *Neurocomputing*, 2025.
- [9] S. Hochreiter and J. Schmidhuber, "Long short-term memory," *Neural Comput.*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [10] W. Hua, Y. Zhou, C. De Sa, Z. Zhang, and G. E. Suh, "Channel gating neural networks," in *Proc. NeurIPS*, 2019.
- [11] N. Shazeer *et al.*, "Outrageously large neural networks: The sparsely-gated mixture-of-experts layer," in *Proc. ICLR*, 2017.
- [12] M. T. Ribeiro, S. Singh, and C. Guestrin, "Why should I trust you? Explaining the predictions of any classifier," in *Proc. KDD*, 2016, pp. 1135–1144.
- [13] S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," in *Advances in NeurIPS*, 2017, pp. 4765–4774.
- [14] M. Sundararajan, A. Taly, and Q. Yan, "Axiomatic attribution for deep networks," in *Proc. ICML*, 2017, pp. 3319–3328.
- [15] D. Bahdanau, K. Cho, and Y. Bengio, "Neural machine translation by jointly learning to align and translate," in *Proc. ICLR*, 2015.
- [16] S. Jain and B. C. Wallace, "Attention is not explanation," in *Proc. NAACL-HLT*, 2019, pp. 3543–3556.
- [17] P. W. Koh *et al.*, "Concept bottleneck models," in *Proc. ICML*, 2020.
- [18] D. P. Kingma and M. Welling, "Auto-encoding variational Bayes," in *Proc. ICLR*, 2014.
- [19] Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner, "Gradient-based learning applied to document recognition," *Proc. IEEE*, vol. 86, no. 11, pp. 2278–2324, 1998.
- [20] A. Krizhevsky, "Learning multiple layers of features from tiny images," *Tech. Rep.*, 2009.
- [21] I. Sharafaldin, A. H. Lashkari, and A. A. Ghorbani, "Toward generating a new intrusion detection dataset and intrusion traffic characterization," in *ICISSP*, 2018, pp. 108–116.